

# **Indexing**

## **Part 2**

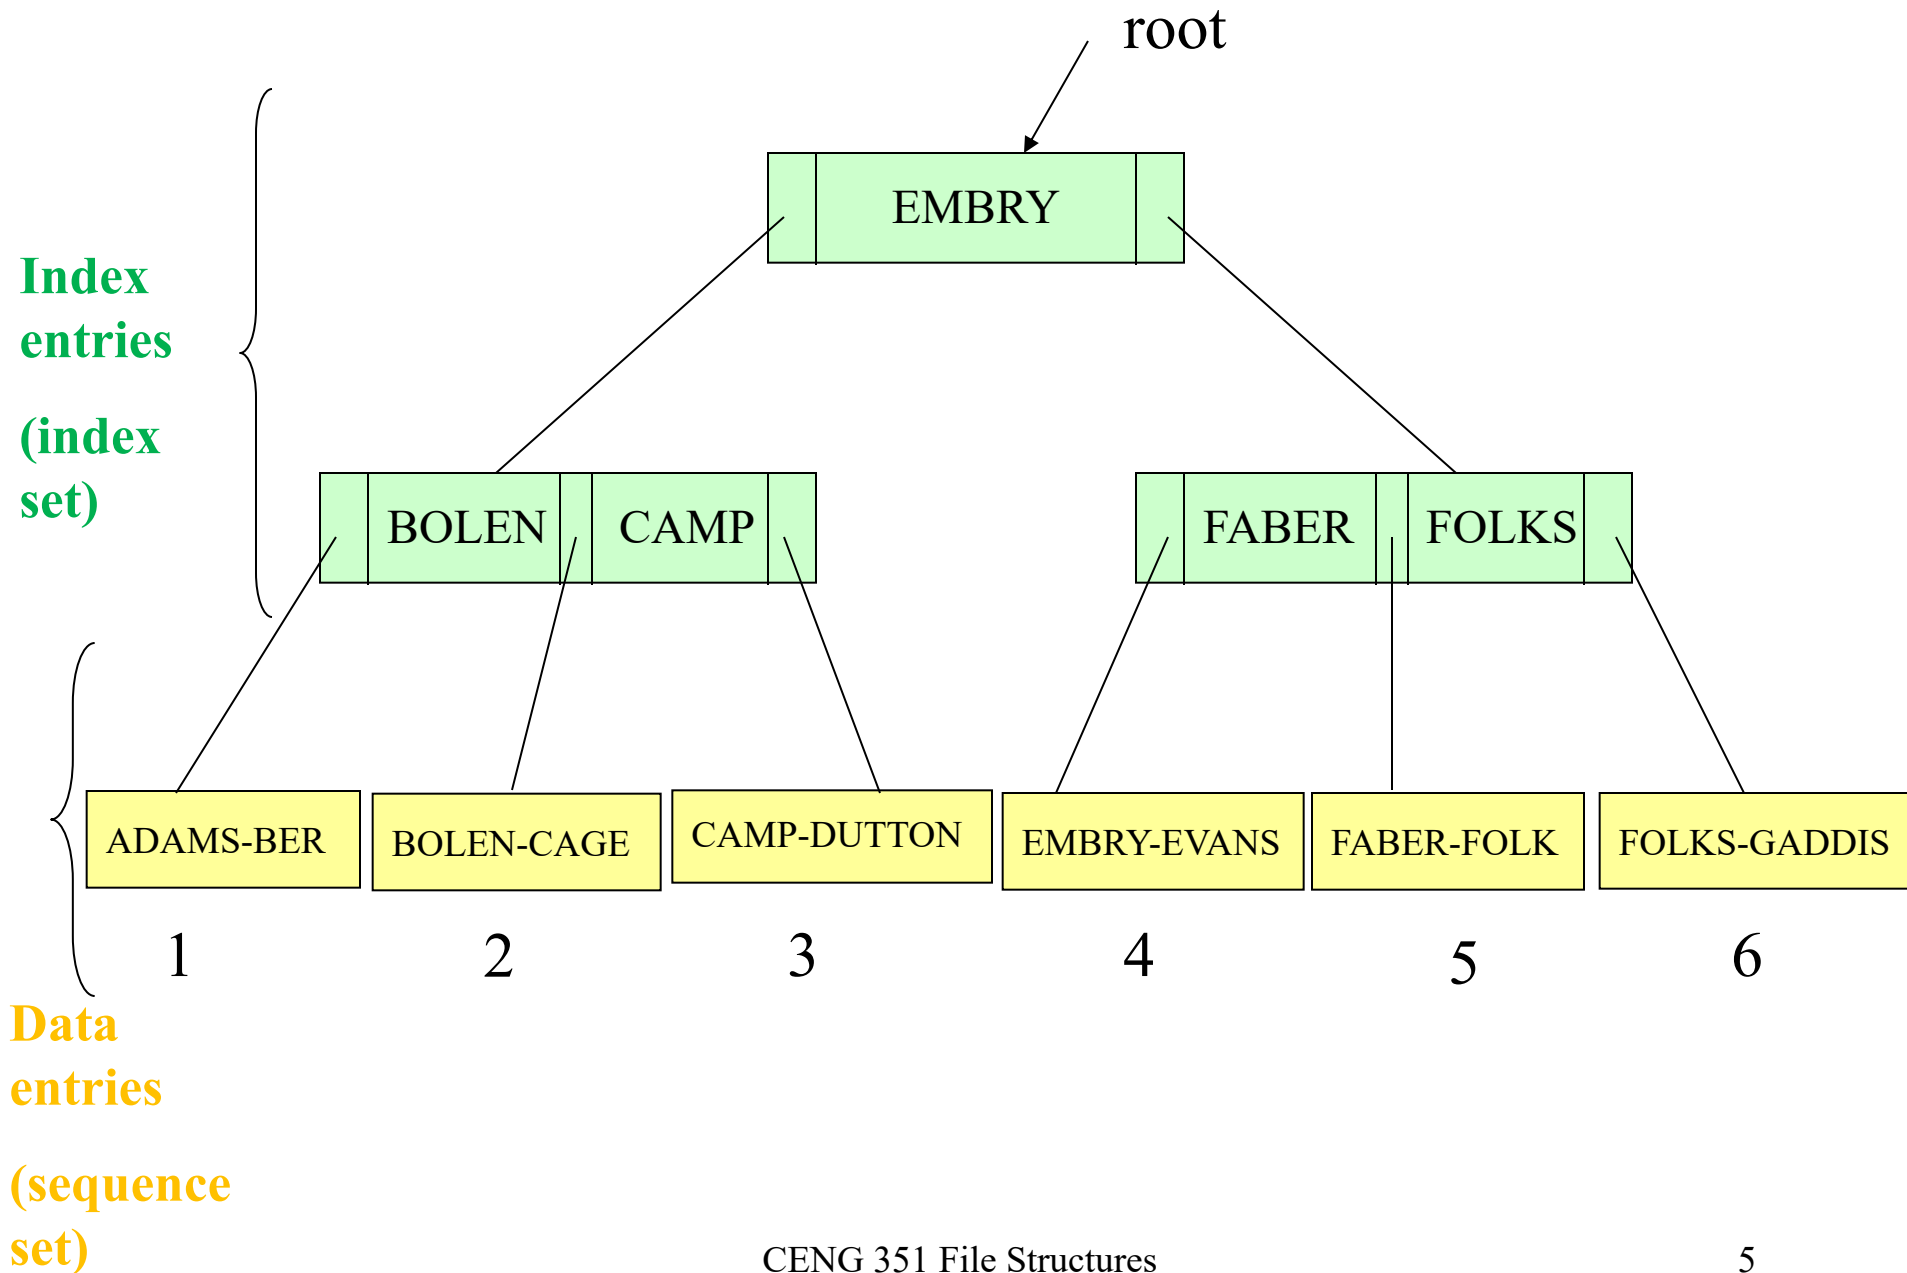
# Tree indexes

- If index doesn't fit in memory:
  - Divide the index structure into blocks,
  - Organize these blocks similarly building a tree structure.
- Tree indexes:
  - B Trees
  - B+ Trees
  - Simple prefix B+ Trees
  - ...

# B+ Trees

- B-tree is one of the most important data structures in computer science.
- What does B stand for? (Not binary!)
- B-tree is a **multiway search** tree.
- Several versions of B-trees have been proposed, but only B+ Trees have been used with large files.
- A B+tree is a B-tree in which data records are in leaf nodes, and faster sequential access is possible.







# Formal definition of B+ Tree Properties

- Properties of a **B+ Tree of order  $d$** :
  - All internal nodes (except root) have **at least  $d$  keys** and **at most  $2d$  keys**.
  - Root can have at least **1 key** and **at most  $2d$  keys**.
  - An internal node with  **$n$  keys** has  **$n+1$  children**
  - The root has at least 2 children unless it's a leaf.
  - All leaves are on the same level (balanced tree).

# B+ tree: Internal/root node structure



# B+ tree: Internal/root node structure




Each  $P_i$  is a pointer to a child node; each  $K_i$  is a search key value  
# of search key values =  $n$ , # of pointers =  $n+1$

In a B+ Tree of order  $d$ :

- All internal nodes (except root) have **at least  $d$  keys** and **at most  $2d$  keys** ( $d \leq n \leq 2d$ ).
- Root can have **at least 1 key** and **at most  $2d$  keys**. ( $1 \leq n \leq 2d$ ).
- An internal node with  **$n$  keys** has  **$n+1$  children**.
- The root has at least 2 children unless it's a leaf.




# B+ tree: Internal/root node structure




Each  $P_i$  is a pointer to a child node; each  $K_i$  is a search key value  
# of search key values =  $n$ ,                      # of pointers =  $n+1$

- Requirements:
- $K_1 < K_2 < \dots < K_n$
- For any search key value  $K$  in the subtree pointed by  $P_i$ ,
  - If  $P_i = P_0$ , we require  $K < K_1$
  - If  $P_i = P_n$ ,  $K_n \leq K$
  - If  $P_i = P_1, \dots, P_{n-1}$ ,  $K_i \leq K < K_{i+1}$

# B+ tree: leaf node structure



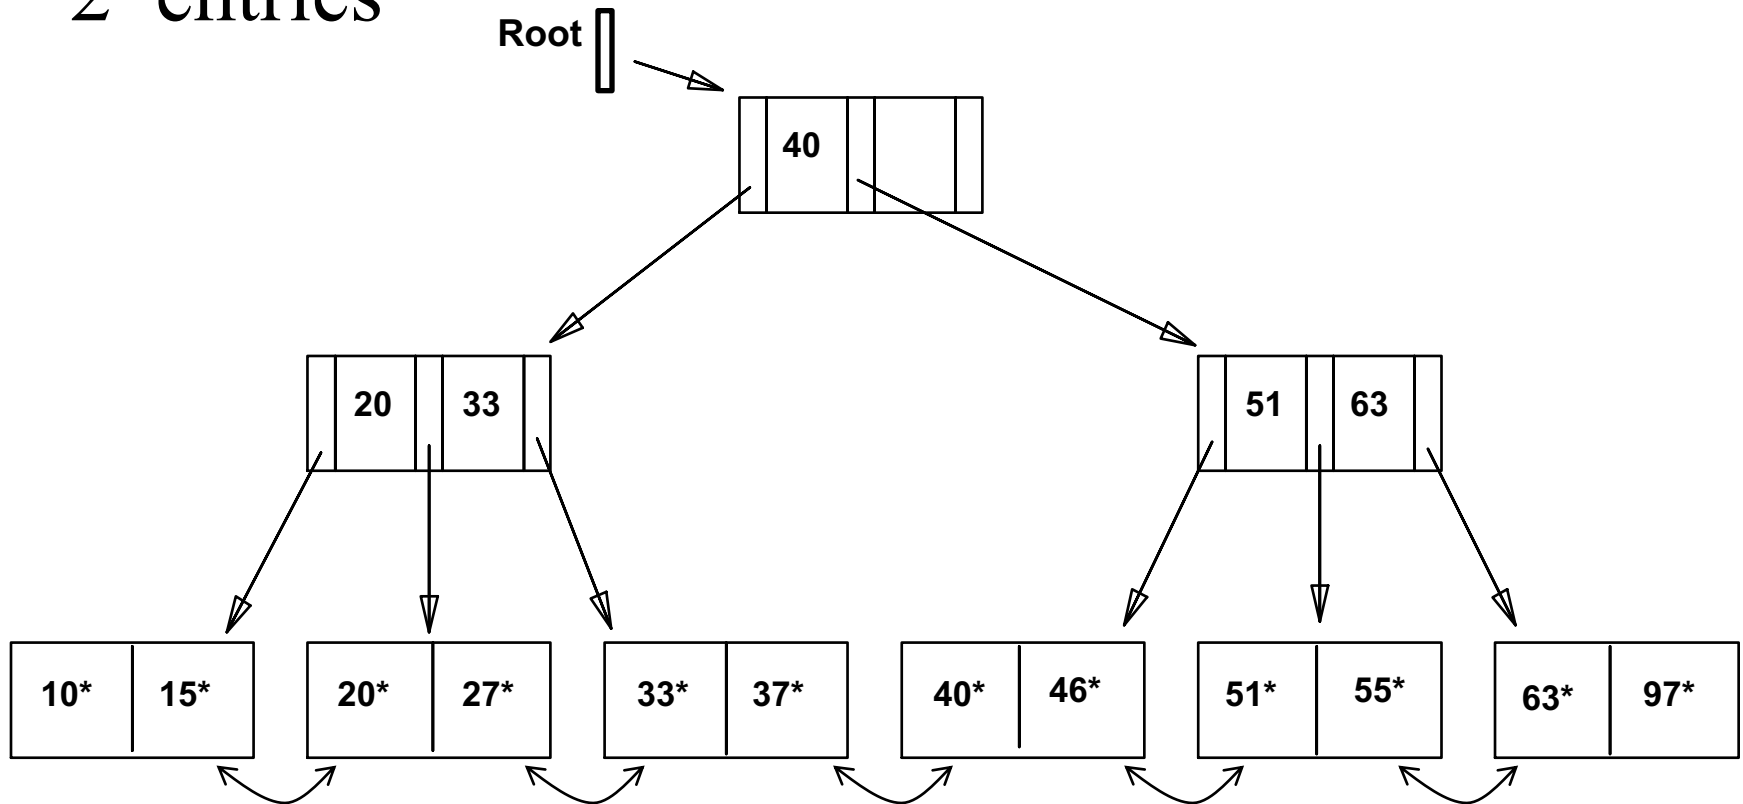
# B+ tree: leaf node structure



- Pointer L points to the left neighbor; R points to the right neighbor (**doubly linked list**)
- $K_1 < K_2 < \dots < K_n$
- $d \leq n \leq 2d$  (d is the order of this B+ tree)
- We will use  $K_i^*$  for the pair  $\langle K_i, r_i \rangle$  and omit L and R for simplicity
- All leaves are on the same level (balanced tree).


# Example: B+ tree with order of 1

- Each node must hold at least 1 entry, and at most 2 entries



# Example: Search in a B+ tree order 2

- Search: how to find the records with a given search key value?
  - Begin at root, and use key comparisons to go to leaf
- Examples: search for 5\*, 16\*, all data entries  $\geq 24^*$  ...
  - The last one is a range search, we need to do the sequential scan, starting from the first leaf containing a value  $\geq 24$ .




# How to Insert a Data Entry into a B+ Tree?

- Let's look at several examples first.

# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root →



Leaf nodes:


$$d \leq n \leq 2d$$

$$2 \leq n \leq 4$$

# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root



**Leaf node overflows!!!**

Leaf nodes:

$$d \leq n \leq 2d$$


$$2 \leq n \leq 4$$



# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root





**Leaf node overflows!!!**

**When a leaf node overflows:**

1) Split the node


First **d** entries stay in old node , move rest of entries to new node



# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root




**Leaf node overflows!!!**

**When a leaf node overflows:**

1) Split the node


First **d** entries stay in old node, move rest of entries to new node




# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root



Leaf node overflows!!!




One new child (leaf node)  
generated; must add **one more  
pointer** to its parent, thus **one  
more key value** as well.

# Inserting 16\*, 8\* into Example B+ tree

$d = 2$

Root



**Leaf node overflows!!!**

**When a leaf node overflows:**


1) Split the node

First  $d$  entries stay, move rest to new node

2) We need a pointer to the new block for the search: **COPY UP** the *middle key* as the search key. Also, add pointer to the new block

# Inserting 16\*, 8\* into Example B+ tree

$d = 2$



**Leaf node overflows!!!**

**When a leaf node overflows:**

1) Split the node


First  $d$  entries stay, move rest to new node

2) **COPY UP** the *middle key* as the search key. Also, add pointer to the new block

# Inserting 8\* (cont.)


$d = 2$

- **Copy up** the middle value (leaf split)



# Inserting 8\* (cont.)

$d = 2$



**Non-leaf node overflows!!!**



**When a non-leaf node overflows:**

1) Split the node

First  $d$  keys  
(and  $d+1$  pointers)  
stay in old node

Last  $d$  keys  
(and  $d+1$  pointers)  
move to new node


2) **PUSH UP** middle  
key (17) and (pointers  
to the blocks)!



# Insertion into B+ tree (cont.)

Recall: Root can have at least **1 key** and **at most  $2d$**  keys.


- Understand difference between **copy-up** and **push-up**
- Observe how min number of entries ( $d$ ) is guaranteed in both leaf and index node splits.



Note that 17 is pushed up and only **appears once** in the index. (Contrast this with a leaf split.)



# Example B+ Tree After Inserting 8\*



Notice that root was split, leading to increase in height.

B+ trees grow **bottom-up** dynamically!

# Inserting a Data Entry into a B+ Tree:

## Summary

- Find correct leaf  $L$ .
- Put data entry onto  $L$ .
  - If  $L$  has enough space, *done!*
  - Else, must split  $L$  (into  $L$  and a new node  $L2$ )
    - Redistribute entries evenly, put middle key in  $L2$
    - copy up middle key.
    - Insert index entry pointing to  $L2$  into parent of  $L$ .
- This can happen recursively
  - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
  - Tree growth: gets wider or one level taller at top.


# Deleting a Data Entry from a B+ Tree

- Examine examples first ...

# Delete 19\* and 20\*

$d = 2$

Root




Leaf nodes:

$$d \leq n \leq 2d$$

$$2 \leq n \leq 4$$

# Delete 19\* and 20\*



$d = 2$




# Delete 19\* and 20\*

$d = 2$

Root



**Delete 19: No problem!**



**Delete 20:  
Leaf node underflows!**


Leaf nodes:

$$d \leq n \leq 2d$$

$$2 \leq n \leq 4$$

# Delete 19\* and 20\*

$d = 2$



**Leaf node underflows!**


**When a leaf node underflows:**

Two options (try in order):

- 1- Redistribute among sibling nodes evenly, and if this is not possible,
- 2- Merge nodes

# Delete 19\* and 20\*

$d = 2$




## Redistribute for leaf nodes:

A **sibling** of a node is the node that is adjacent (immediate to left or right) to it, and has the same parent



# Delete 19\* and 20\*

$d = 2$




## Redistribute for leaf nodes:

A **sibling** of a node is the node that is adjacent (immediate to left or right) to it, and has the same parent

# Delete 19\* and 20\*

$d = 2$




## Redistribute for leaf nodes:

A **sibling** of a node is the node that is adjacent (immediate to left or right) to it, and has the same parent

# Delete 19\* and 20\*

$d = 2$



**Sibling & has > d entries!**


## Redistribute for leaf nodes:

A **sibling** of a node is the node that is adjacent (immediate to left or right) to it, and has the same parent

1) Redistribute among siblings

# Delete 19\* and 20\*

$d = 2$




**Sibling & has  $> d$  entries!**

## Redistribute for leaf nodes:

A **sibling** of a node is the node that is adjacent (immediate to left or right) to it, and has the same parent

- 1) Redistribute among siblings
- 2) COPY-UP (Update) the middle key as the search key

# Deleting 19\* and 20\* (cont.)




- Notice how 27 is *copied up*.
- But can we move it up?
- Now we want to delete 24
- Underflow again!

# Delete 24\*

$d = 2$

Root



**Option (1) Not Applicable here!**

**When a leaf node underflows:**

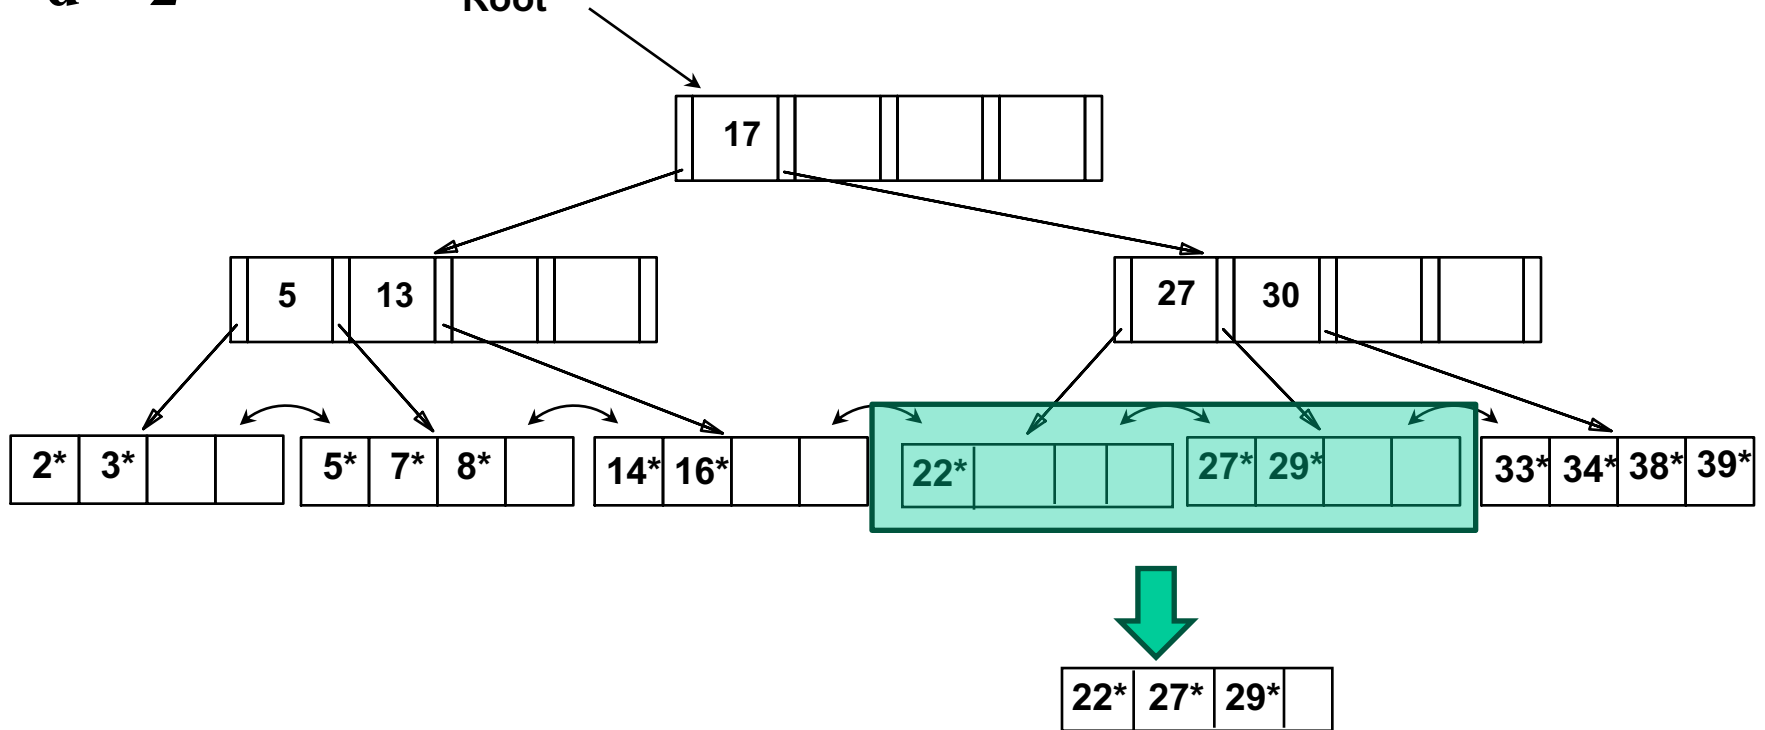
Two options (try in order):

- 1- Redistribute among sibling nodes evenly, and if this is not possible,
- 2- Merge nodes

# Delete 24\*

$d = 2$

Root




**Option (2): Merge leaf nodes**

Step 1: Merge leaf nodes

# Delete 24\*

$d = 2$

Root



**Option (2): Merge leaf nodes**

Step 1: Merge leaf nodes


Step 2: Remove the search key entry and pointer to the discarded node



# Delete 24\*

$d = 2$

Root



Is it good like this?

**Option (2): Merge leaf nodes**


Step 1: Merge leaf nodes

Step 2: Remove the search key entry and pointer to the discarded node

# Delete 24\*

$d = 2$


Root



# Delete 24\*

$d = 2$

Root



## When a non-leaf node underflows:


Two options (try in order):

- 1- Redistribute among sibling nodes evenly evenly, and if this is not possible,
- 2- Merge nodes

# Delete 24\*

$d = 2$

Root



**Option (1) Not Applicable here!**


**When a non-leaf node underflows:**

Two options (try in order):

- 1- Redistribute among sibling nodes evenly evenly, and if this is not possible,
- 2- Merge nodes

# Delete 24\*

d = 2



## Option (2): Merge non-leaf nodes

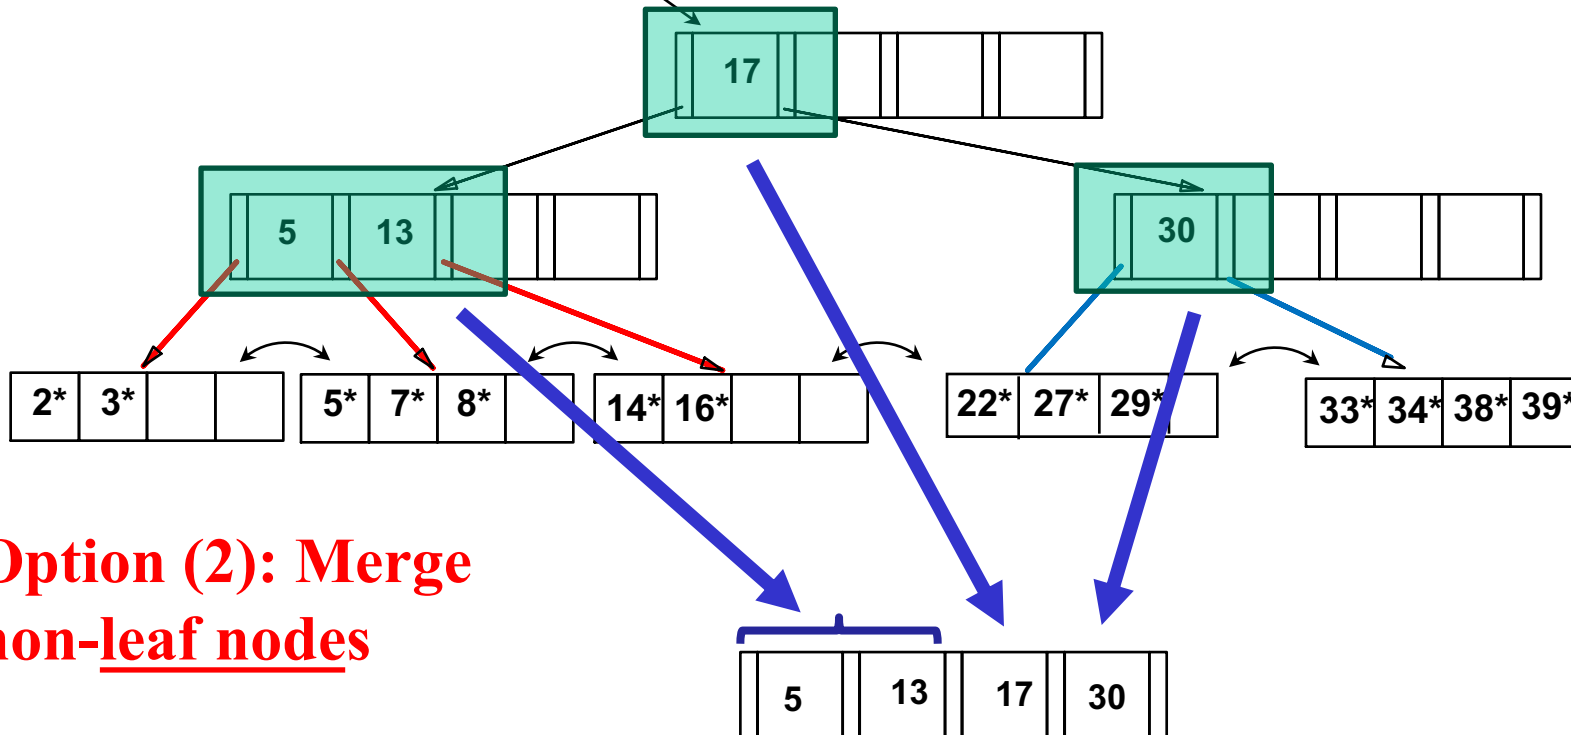
Merge:

- Entries in first non-leaf node (together with pointers),
- **PULL DOWN** the splitting search key,
- followed by the entries in the second non-leaf node (together with pointers)

$d = 2$

Root

# Delete 24\*




**Option (2): Merge  
non-leaf nodes**

Merge:

- Entries in first non-leaf node (together with pointers),
- **PULL DOWN** the splitting search key,
- followed by the entries in the second non-leaf node (together with pointers)

# Delete 24\*



# Deleting a Data Entry from a B+ Tree:

## Summary

- Start at root, find leaf  $L$  where entry belongs.
- Remove the entry.
  - If  $L$  is at least half-full, *done!*
  - If  $L$  has only **d-1** entries,
    - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as  $L$* ).
    - If re-distribution fails, merge  $L$  and sibling.
- If merge occurred, must delete entry (pointing to  $L$  or sibling) from parent of  $L$ .
- Merge could propagate to root, decreasing height.



# Non-leaf Node Redistribution

**When a non-leaf node underflows:**

Two options (try in order):


- 1- Redistribute among nodes evenly, and if this is not possible,
- 2- Merge nodes

We have already seen an example for the second case.

How about the first case!

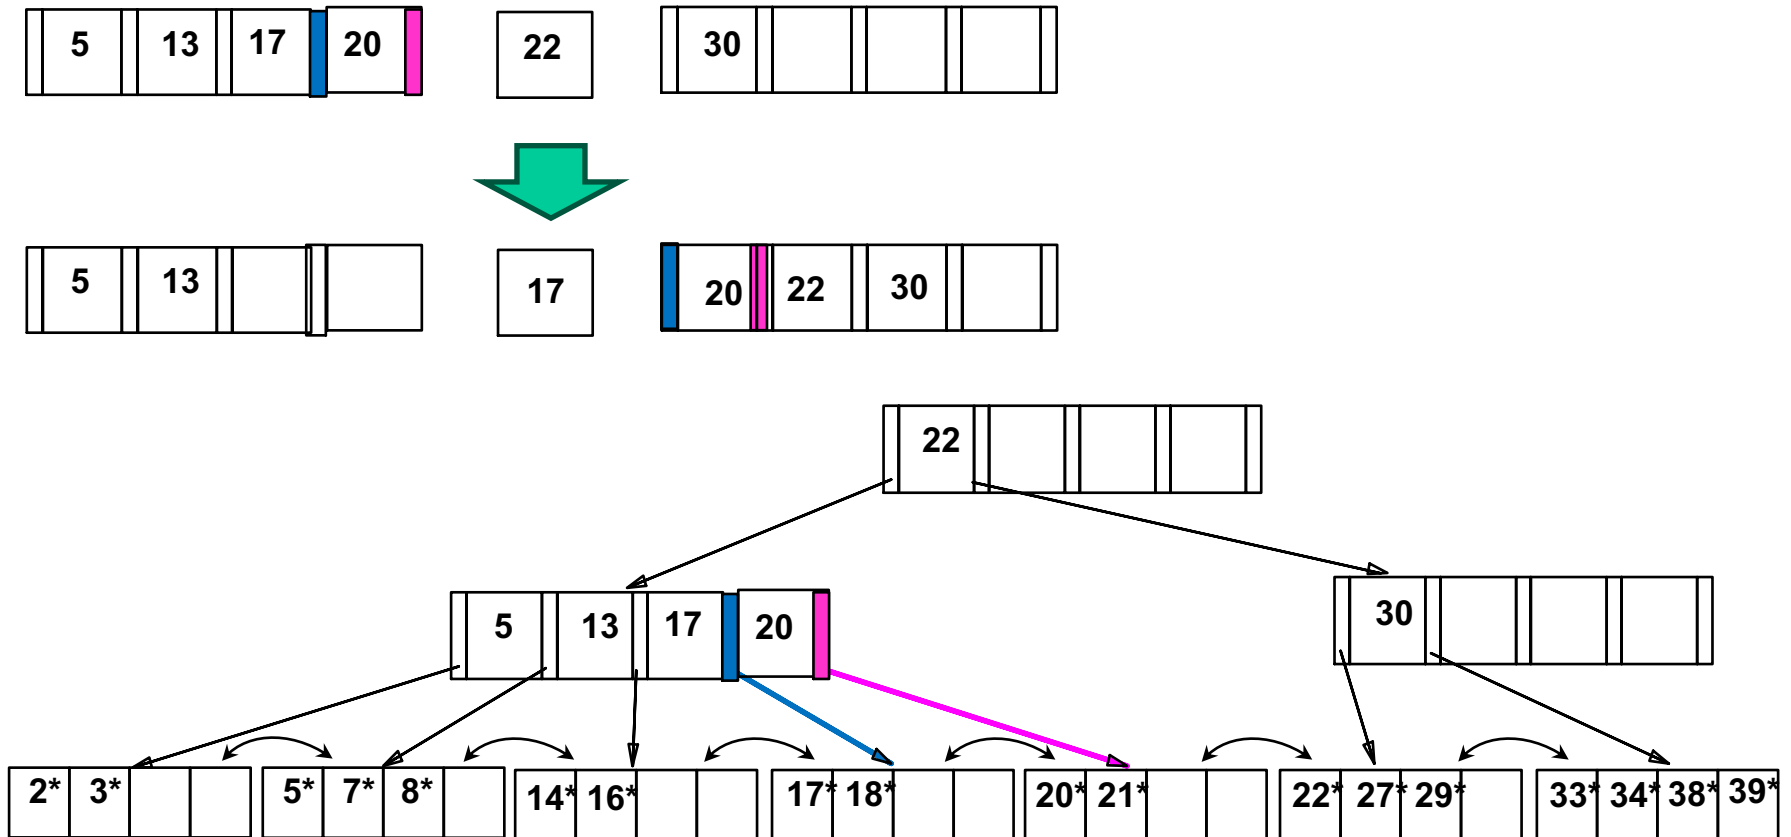
# Example of Non-leaf Re-distribution

- Tree is shown below *during deletion* of 24\*. (What could be a possible initial tree?)
- In contrast to previous example, can re-distribute entry from left child of root to right child.




# Re-distribution

- Intuitively, entries are *re-distributed by 'pushing through'* the splitting entry in the parent node.
- Consider all search keys together




# After Re-distribution



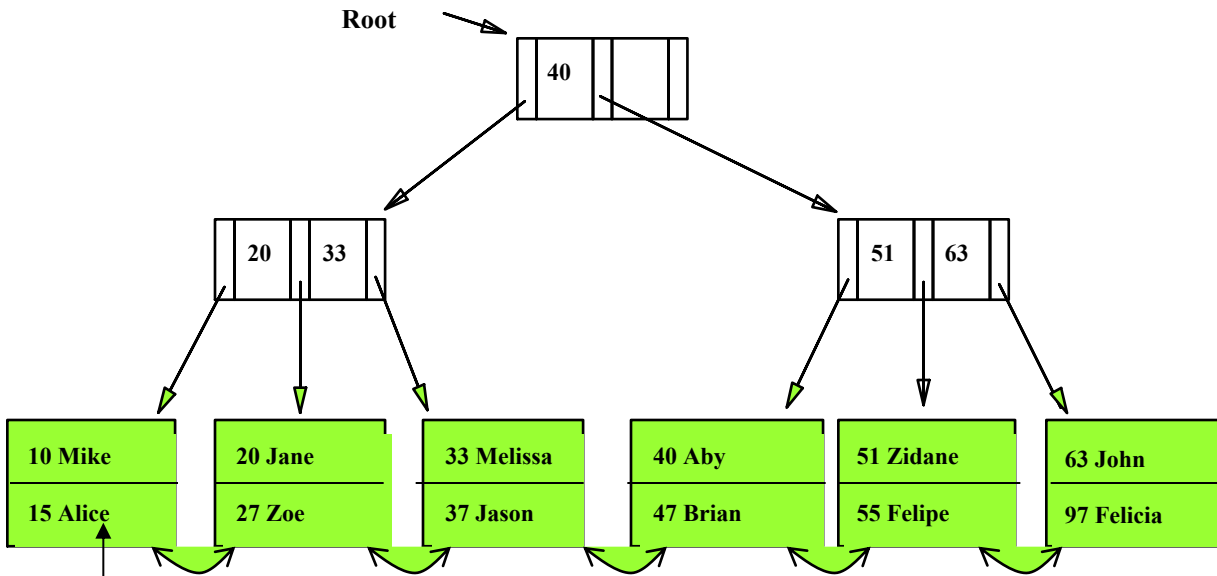
# Primary vs Secondary Index

- Note: We were assuming the data items were in sorted order
  - This is called *primary/clustered B+tree* index
- *Secondary B+tree* index:
  - Built on an attribute that the file is not sorted on.
- Can have many different indexes on the same file.

# A Secondary B+-Tree index



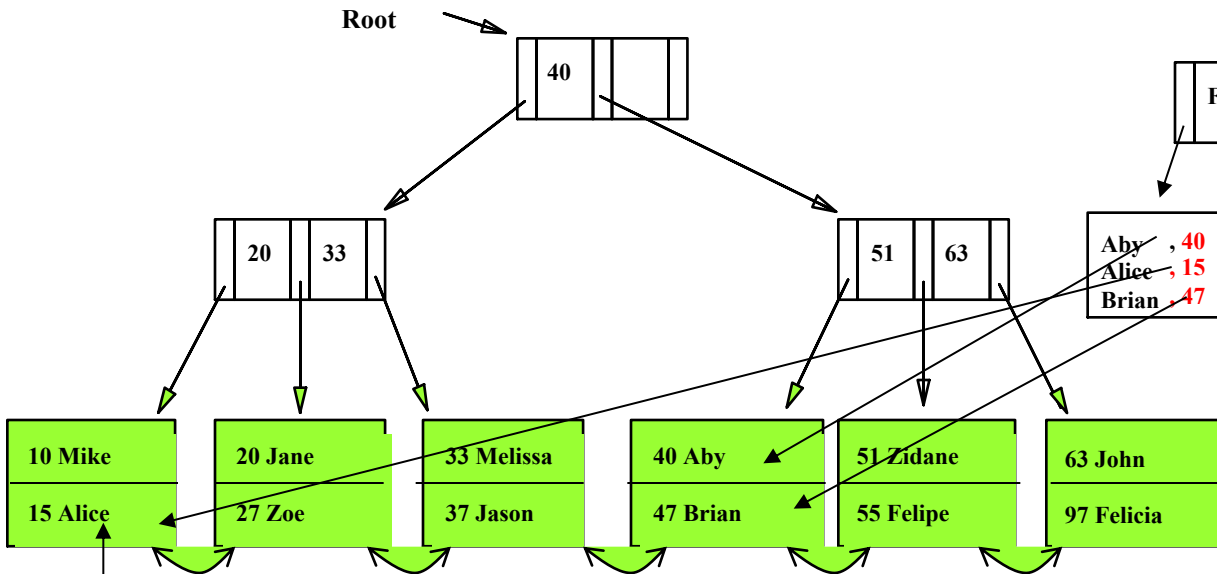
A file organized as (or, has) a  
**Primary B+-Tree** index on *ssn*




- As 15\*, we store the **actual data record** with key value 15 (**Alternative-1**)
- In this case, the leaf nodes can be larger, say a few blocks (pages)

A file organized as (or, has) a **Primary B+-Tree** index on *ssn*

The same file also has a **Secondary B+-Tree** index on *name*



- As 15\*, we store the **actual data record** with key value 15 (**Alternative-1**)
- In this case, the leaf nodes can be larger, say a few blocks (pages)




- We have  $k^*$  as  $\langle \text{key}, \text{rid} \rangle$  (**Alternative 2**)
- We can have rid's as pointers, or use PK as rid. We show both above



# Cost for searching a value in B+ tree

- Assumptions:
  - Each interior node is a disk block
  - Each leaf node is also a disk block and data entries ( $K^*$ ) are of the form  $\langle \text{key}, \text{ptr} \rangle$ . There are  $D$  data entries.
  - Let  $F$  be the average number of pointers in a node (for internal nodes, it is called *fanout*, i.e., avg. number of children )
- Observe: Let  $H$  be the height of the B+ tree: we need to read  $H+1$  nodes (blocks) to reach a data entry in a leaf node
- How do we find  $H$ ?
  - Level 1 = 1 page =  $F^0$  page
  - Level 2 =  $F$  pages =  $F^1$  pages
  - Level 3 =  $F * F$  pages =  $F^2$  pages
  - Level  $H+1$  = ..... =  $F^H$  pages (i.e., leaf nodes)
  - $F$  pointers  $\rightarrow F-1$  keys, so there must be  $D/(F-1)$  leaf nodes
  - $D/(F-1) = F^H$ . That is,  $H = \log_F(\frac{D}{F-1})$



# B+ Trees in Practice

- Typical order: 100. Typical fill-factor: 66%.
  - average fanout = 133 (i.e, # of pointers in internal node)
- Can often hold top levels in buffer pool:
  - Level 1 = 1 page = 8 Kbytes
  - Level 2 = 133 pages = 1 Mbyte
  - Level 3 = 17,689 pages = 133 MBytes
- Suppose there are 1,000,000,000 data entries.
  - $H = \log_{133}(1000000000/132) < 4$
  - The cost is reading  $H+1 = 5$  pages

# Cost Computation: Another Example

Leaves would store the **actual records**

- A primary B+ tree index on key field giftID.
- 2.500.000 gift records, each record: 400 bytes.
- giftID: 12 bytes, address pointer: 4 bytes
- A bucket can hold 500 records
  - So we have larger leaf nodes (called **buckets**), as we store actual records
  - No claim for interior nodes, assume each is a block!
- B+ tree will have a fill factor of 50% [min occupancy]
- B (block size): 1600
- s: 10 ms, r: 5 ms, btt: 1 ms.

## **a) No of index nodes and their total size**

We need to find i) fanout of the nodes, and ii) no of leaves.

**i) fanout:** Assume , **n keys (n+1) ptrs** can fit to an index node:

$$n \times 12 + (n+1) \times 4 = 1600 \text{ bytes} \rightarrow 16n = 1596 / 16 \rightarrow n = 99$$

So at most 99 keys in a node (  $2d = 99$ ,  $d$  (tree order) is  $\text{floor}(99/2)$  )

Tree fill factor 50%; max 99 keys  $\times$  50% = 49 keys

fanout:  $49 + 1 = 50$  ptrs per node

**ii) no of leaves:**

$$500 \text{ rec/leaf} \times \text{fill factor (50\%)} = 250 \text{ recs/leaf}$$

$$2.5\text{M records} / 250 = 10000 \text{ leaf nodes (i.e., buckets)}$$

## **a) No of index nodes and their total size**

- Tree height =  $\log_{50} 10000 = 3$
- So, there are  $H+1 = 4$  levels

Level 4: 10000 leaf nodes (data buckets)

Level 3:  $\text{ceil}(10000 / 50 \text{ ptrs}) = 200$  nodes

Level 2:  $\text{ceil}(200/50) = 4$  nodes

Level 1:  $\text{ceil}(4/50) = 1$  node (root)

Index nodes:  $1 + 4 + 200 = 205$

Total Size:  $205 \times 1600$  bytes

## **b) Time cost of reading an arbitrary record**

- Three has  $H=3$ , so 4 levels
- At the first 3 levels, we fetch index nodes:  
 $3 \times (s + r + btt) = 3 \times (10 + 5 + 1) = 48 \text{ ms}$
- At the fourth level we fetch the leaf node (data bucket)
  - But how many blocks is a data bucket?
  - $(500 \text{ recs} \times 400 \text{ bytes/rec}) / 1600 = 125 \text{ blocks}$
  - So, cost  $s + r + 125 \times btt = 10 + 5 + 125 \times 1 = 140 \text{ ms}$
- Total cost:  $48 + 140 = 188 \text{ ms}$

## c) Cost of reading all records in sorted manner

- Reach to leftmost leaf node, as before:
- at the first 3 levels, we fetch index nodes:  
$$3 \times (s + r + btt) = 3 \times (10 + 5 + 1) = 48 \text{ ms}$$
- Read all the leaf nodes (using doubly linked list pointers)
  - $10000 (s + r + 125 \times btt)$
- Think: What if this is a secondary B+ tree and we store  $\langle \text{key}, \text{ptr} \rangle$  pairs at leaf nodes (data buckets)?


# Terminology

- **Blocking Factor:** the number of records which can fit in a leaf node.
- **Fan-out :** the average number of children of an internal node.
- A B+tree index can be used either as a primary index or a secondary index.
  - **Primary index:** determines the way the records are actually stored
  - **Secondary index:** the records in the file are not grouped in blocks according to keys of secondary indexes



# Bulk Loading of a B+ Tree


- If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
- Bulk Loading can be done much more efficiently.
  - Initialization: Sort all data entries, insert pointer to first (leaf) page in a new (root) page



# Bulk Loading (Contd.)

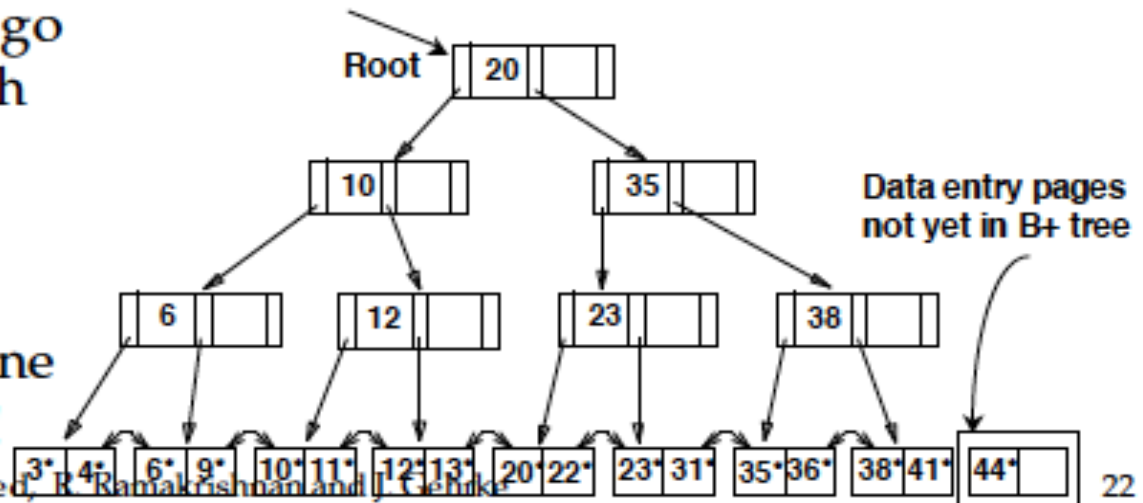


- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level.



When this fills up, it splits. (Split may go up right-most path to the root.)

- ❖ Much faster than repeated inserts, especially when one considers locking!



# Summary

- Tree-structured indexes are ideal for range-searches, also good for equality searches.
- B+ tree is a dynamic structure.
  - Inserts/deletes leave tree height-balanced; High fanout (**F**) means depth rarely more than 3 or 4.
  - Almost always better than maintaining a sorted file.
  - Typically, 67% occupancy on average.
  - If data entries are data records, splits can change rids!
- Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.

# More...

- Hash-based Indexes
  - Static Hashing
  - Extendible Hashing
  - Linear Hashing
- Grid-files
- R-Trees
- etc...
- A nice animation site for B+ trees:  
<https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>